



**Machine Vision ( CSE480)**

**Lab 4 Report**


Name	ID	Section
Mahmoud Elsayd Abdelqader Labib Eldwakhly	21P0017	1

Submitted to Dr Hossam Hassan & Eng. Dina Zakaria

Fall 2025

## Task 1 :

### Dataset Exploration :




```
==== MODEL SUMMARY ====
Model: "sequential_1"
```


Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 128)	100,480
dropout_1 (Dropout)	(None, 128)	0
dense_4 (Dense)	(None, 64)	8,256
dense_5 (Dense)	(None, 10)	650

```
Total params: 109,386 (427.29 KB)
Trainable params: 109,386 (427.29 KB)
Non-trainable params: 0 (0.00 B)
```

```
model = tf.keras.Sequential([
    tf.keras.layers.Input(shape=(784,)),
    tf.keras.layers.Dense(128, activation="relu"),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(64, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")])
```



Random Test Cases: True vs Predicted + Confidence



#### Model:

- Dense-only network: 128(ReLU) + Dropout + 64(ReLU) + 10(Softmax).

#### Training (10 epochs):

- Final Train Loss: 0.0500 | Final Val Loss: 0.0683
- Final Train Acc : 0.9836 | Final Val Acc : 0.9817

#### Testing:

- Test Loss: 0.0697 | Test Accuracy: 0.9780

```

import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

tf.random.set_seed(42)
np.random.seed(42)

# -----
# 1) Load MNIST
# -----
(x_train_img, y_train), (x_test_img, y_test) = tf.keras.datasets.mnist.load_data()

print("== DATASET OVERVIEW ==")
print("Train images shape:", x_train_img.shape, " Train labels shape:", y_train.shape)
print("Test images shape:", x_test_img.shape, " Test labels shape:", y_test.shape)
print("Pixel range (train):", (x_train_img.min(), x_train_img.max()))
print("Pixel range (test): ", (x_test_img.min(), x_test_img.max()))
print("Classes:", np.unique(y_train))

# -----
# 2) Explore dataset
# -----
def plot_class_distribution(labels, title):
    counts = np.bincount(labels, minlength=10)
    plt.figure(figsize=(8,3))
    plt.bar(range(10), counts)
    plt.xticks(range(10))
    plt.xlabel("Digit class")
    plt.ylabel("Count")
    plt.title(title)
    plt.show()
    return counts

train_counts = plot_class_distribution(y_train, "Train Class Distribution")
test_counts = plot_class_distribution(y_test, "Test Class Distribution")

print("Train class counts:", train_counts)
print("Test class counts:", test_counts)

def show_samples(images, labels, n=12, title="Random Samples"):
    idx = np.random.choice(len(images), n, replace=False)
    cols = 6
    rows = int(np.ceil(n / cols))
    plt.figure(figsize=(12, 2*rows))
    for i, k in enumerate(idx):
        plt.subplot(rows, cols, i+1)
        plt.imshow(images[k], cmap="gray")
        plt.title(f"y={labels[k]}")
        plt.axis("off")
    plt.suptitle(title)
    plt.tight_layout()
    plt.show()

show_samples(x_train_img, y_train, n=12, title="Random Training Samples")

# -----
# 3) Preprocess: normalize + flatten
# -----
x_train = (x_train_img.astype("float32") / 255.0).reshape(-1, 28*28)
x_test = (x_test_img.astype("float32") / 255.0).reshape(-1, 28*28)

# -----
# 4) Build Dense-only model
# -----
model = tf.keras.Sequential([
    tf.keras.layers.Input(shape=(784,)),

```

```

        tf.keras.layers.Dense(128, activation="relu"),
        tf.keras.layers.Dropout(0.2),
        tf.keras.layers.Dense(64, activation="relu"),
        tf.keras.layers.Dense(10, activation="softmax")
    )

model.compile(
    optimizer=tf.keras.optimizers.Adam(1e-3),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

print("\n==== MODEL SUMMARY ====")
model.summary()

# -----
# 5) Train for 10 epochs
# -----
EPOCHS = 10
history = model.fit(
    x_train, y_train,
    validation_split=0.1,
    epochs=EPOCHS,
    batch_size=128,
    verbose=1
)

# -----
# 6) Plot training curves
# -----
def plot_training_curves(hist):
    h = hist.history
    ep = range(1, len(h["loss"]) + 1)
    plt.figure(figsize=(12,4))

    plt.subplot(1,2,1)
    plt.plot(ep, h["loss"], label="Train Loss")
    plt.plot(ep, h["val_loss"], label="Val Loss")
    plt.xlabel("Epoch"); plt.ylabel("Loss")
    plt.title("Loss Curves"); plt.legend()

    plt.subplot(1,2,2)
    plt.plot(ep, h["accuracy"], label="Train Acc")
    plt.plot(ep, h["val_accuracy"], label="Val Acc")
    plt.xlabel("Epoch"); plt.ylabel("Accuracy")
    plt.title("Accuracy Curves"); plt.legend()

    plt.tight_layout()
    plt.show()

plot_training_curves(history)

# -----
# 7) Evaluate on test set
# -----
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)
print("\n==== TEST PERFORMANCE ====")
print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_acc:.4f}")

# -----
# 8) Test cases (predictions) with output
# -----
def predict_test_cases(model, x_test_flat, x_test_images, y_test, n_cases=12):
    idx = np.random.choice(len(x_test_flat), n_cases, replace=False)
    probs = model.predict(x_test_flat[idx], verbose=0)

```

```

preds = np.argmax(probs, axis=1)
confs = np.max(probs, axis=1)

cols = 6
rows = int(np.ceil(n_cases / cols))
plt.figure(figsize=(12, 2*rows))
for i, k in enumerate(idx):
    plt.subplot(rows, cols, i+1)
    plt.imshow(x_test_images[k], cmap="gray")
    correct = (preds[i] == y_test[k])
    mark = "✓" if correct else "✗"
    plt.title(f"T={y_test[k]} P={preds[i]} ({confs[i]*100:.1f}%) {mark}")
    plt.axis("off")
plt.suptitle("Random Test Cases: True vs Predicted + Confidence")
plt.tight_layout()
plt.show()

# Also print a small table-like output
print("\nSample Predictions (index, true, pred, confidence):")
for i, k in enumerate(idx):
    print(f"- idx={k:5d} | true={y_test[k]} | pred={preds[i]} | conf={confs[i]:.4f}")

predict_test_cases(model, x_test, x_test_img, y_test, n_cases=12)

# -----
# 9) Confusion Matrix (no sklearn needed)
# -----
def confusion_matrix_np(y_true, y_pred, num_classes=10):
    cm = np.zeros((num_classes, num_classes), dtype=np.int32)
    for t, p in zip(y_true, y_pred):
        cm[t, p] += 1
    return cm

all_probs = model.predict(x_test, verbose=0)
y_pred = np.argmax(all_probs, axis=1)

cm = confusion_matrix_np(y_test, y_pred, num_classes=10)

plt.figure(figsize=(6,5))
plt.imshow(cm, cmap="Blues")
plt.title("Confusion Matrix (Test)")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.xticks(range(10))
plt.yticks(range(10))
plt.colorbar()
plt.tight_layout()
plt.show()

# Print a few most confused pairs
off_diag = cm.copy()
np.fill_diagonal(off_diag, 0)
pairs = []
for i in range(10):
    for j in range(10):
        if i != j and off_diag[i, j] > 0:
            pairs.append((off_diag[i, j], i, j))
pairs.sort(reverse=True)

print("\n== MOST COMMON CONFUSIONS (true -> predicted) ===")
for c, t, p in pairs[:10]:
    print(f"- {t} -> {p}: {c} times")

# -----
# 10) Show misclassified examples
# -----

```

```

def show_misclassified(x_test_images, y_true, y_pred, probs, n=12):
    wrong_idx = np.where(y_true != y_pred) [0]
    if len(wrong_idx) == 0:
        print("No misclassifications found (rare).")
        return
    pick = np.random.choice(wrong_idx, size=min(n, len(wrong_idx)), replace=False)

    cols = 6
    rows = int(np.ceil(len(pick) / cols))
    plt.figure(figsize=(12, 2*rows))
    for i, k in enumerate(pick):
        plt.subplot(rows, cols, i+1)
        plt.imshow(x_test_images[k], cmap="gray")
        conf = np.max(probs[k])
        plt.title(f"T={y_true[k]} P={y_pred[k]} ({conf*100:.1f}%)")
        plt.axis("off")
    plt.suptitle("Misclassified Test Examples")
    plt.tight_layout()
    plt.show()



show_misclassified(x_test_img, y_test, y_pred, all_probs, n=12)

# -----
# 11) Small report
# -----
train_loss_last = history.history["loss"] [-1]
val_loss_last   = history.history["val_loss"] [-1]
train_acc_last  = history.history["accuracy"] [-1]
val_acc_last    = history.history["val_accuracy"] [-1]


print("\n" + "="*60)
print("SMALL DETAILED REPORT")
print("="*60)
print("Exploration:")
print("- MNIST has 60,000 train images and 10,000 test images, each 28x28 grayscale.")
print("- We displayed random samples and plotted class distributions (counts per digit).")
print("\nModel:")
print("- Dense-only network: 128(ReLU) + Dropout + 64(ReLU) + 10(Softmax).")
print("\nTraining (10 epochs):")
print(f"- Final Train Loss: {train_loss_last:.4f} | Final Val Loss: {val_loss_last:.4f}")
print(f"- Final Train Acc : {train_acc_last:.4f} | Final Val Acc : {val_acc_last:.4f}")
print("\nTesting:")
print(f"- Test Loss: {test_loss:.4f} | Test Accuracy: {test_acc:.4f}")
print("\nTest Cases & Errors:")
print("- We predicted random test cases and showed True vs Predicted with confidence.")
print("- We plotted a confusion matrix + printed most common confusions.")
print("- We displayed some misclassified examples to inspect error patterns.")
print("="*60)

```

## Task 2 :



### Random Test Cases: True vs Predicted + Confidence (Fashion-MNIST)

T=Sandal P=Sandal 100.0% ✓	T=T-shirt/top P=T-shirt/top 81.9% ✓	T=T-shirt/top P=T-shirt/top 61.8% ✓	T=Dress P=Coat 41.1% ✗	T=Trouser P=Trouser 100.0% ✓
 T=Ankle boot P=Ankle boot 100.0% ✓	 T=Trouser P=Trouser 98.7% ✓	 T=Pullover P=Shirt 67.1% ✗	 T=T-shirt/top P=T-shirt/top 74.4% ✓	 T=Bag P=Bag 99.5% ✓

#### Preprocessing:

- Normalized pixel values to  $[0, 1]$ .
- Flattened each image from 28x28 to a 784-length vector.

#### Model (Dense-only):

- Dense(256, ReLU) + Dropout(0.3) + Dense(128, ReLU) + Dense(10, Softmax).

#### Training (10 epochs):

- Final Train Loss: 0.2882 | Final Val Loss: 0.3122
- Final Train Acc : 0.8923 | Final Val Acc : 0.8828

#### Testing:

- Test Loss: 0.3378 | Test Accuracy: 0.8794

```

import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

tf.random.set_seed(42)
np.random.seed(42)

# -----
# 1) Load Fashion-MNIST
# -----
(x_train_img, y_train), (x_test_img, y_test) = tf.keras.datasets.fashion_mnist.load_data()

class_names = [
    "T-shirt/top", "Trouser", "Pullover", "Dress", "Coat",
    "Sandal", "Shirt", "Sneaker", "Bag", "Ankle boot"
]

print("==== DATASET OVERVIEW (Fashion-MNIST) ===")
print("Train images shape:", x_train_img.shape, " Train labels shape:", y_train.shape)
print("Test images shape:", x_test_img.shape, " Test labels shape:", y_test.shape)
print("Pixel range (train):", (x_train_img.min(), x_train_img.max()))
print("Pixel range (test): ", (x_test_img.min(), x_test_img.max()))
print("Classes (0-9):", list(range(10)))
print("Class names:", class_names)

# -----
# 2) Explore dataset (distribution + sample images)
# -----
def plot_class_distribution(labels, title):
    counts = np.bincount(labels, minlength=10)
    plt.figure(figsize=(10,3))
    plt.bar(range(10), counts)
    plt.xticks(range(10), class_names, rotation=35, ha="right")
    plt.xlabel("Class")
    plt.ylabel("Count")
    plt.title(title)
    plt.tight_layout()
    plt.show()
    return counts

train_counts = plot_class_distribution(y_train, "Train Class Distribution (Fashion-MNIST)")
test_counts = plot_class_distribution(y_test, "Test Class Distribution (Fashion-MNIST)")

print("Train class counts:", train_counts)
print("Test class counts:", test_counts)

def show_samples(images, labels, n=12, title="Random Samples"):
    idx = np.random.choice(len(images), n, replace=False)
    cols = 6
    rows = int(np.ceil(n / cols))
    plt.figure(figsize=(12, 2.2*rows))
    for i, k in enumerate(idx):
        plt.subplot(rows, cols, i+1)
        plt.imshow(images[k], cmap="gray")
        plt.title(class_names[labels[k]])
        plt.axis("off")
    plt.suptitle(title)
    plt.tight_layout()
    plt.show()

show_samples(x_train_img, y_train, n=12, title="Random Training Samples (Fashion-MNIST)")

# -----
# 3) Normalize + Flatten
# -----
x_train = (x_train_img.astype("float32") / 255.0).reshape(-1, 28*28)
x_test = (x_test_img.astype("float32") / 255.0).reshape(-1, 28*28)

# -----
# 4) Build Dense-only model
# -----
model = tf.keras.Sequential([
    tf.keras.layers.Input(shape=(784,)),
    tf.keras.layers.Dense(256, activation="relu"),
    tf.keras.layers.Dropout(0.3),
    tf.keras.layers.Dense(128, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")
])

model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),
    loss="sparse_categorical_crossentropy",
)

```

```

    metrics=["accuracy"]
)

print("\n==== MODEL SUMMARY ===")
model.summary()

# -----
# 5) Train for 10 epochs
# -----
EPOCHS = 10
history = model.fit(
    x_train, y_train,
    epochs=EPOCHS,
    batch_size=128,
    validation_split=0.1,
    verbose=1
)

# -----
# 6) Plot training curves (loss + accuracy)
# -----
def plot_training_curves(hist):
    h = hist.history
    ep = range(1, len(h["loss"]) + 1)
    plt.figure(figsize=(12,4))

    plt.subplot(1,2,1)
    plt.plot(ep, h["loss"], label="Train Loss")
    plt.plot(ep, h["val_loss"], label="Val Loss")
    plt.xlabel("Epoch"); plt.ylabel("Loss")
    plt.title("Loss Curves"); plt.legend()

    plt.subplot(1,2,2)
    plt.plot(ep, h["accuracy"], label="Train Acc")
    plt.plot(ep, h["val_accuracy"], label="Val Acc")
    plt.xlabel("Epoch"); plt.ylabel("Accuracy")
    plt.title("Accuracy Curves"); plt.legend()

    plt.tight_layout()
    plt.show()

plot_training_curves(history)

# -----
# 7) Test evaluation
# -----
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)
print("\n==== TEST PERFORMANCE ===")
print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_acc:.4f}")

# -----
# 8) Predictions + Confusion Matrix (colored)
# -----
all_probs = model.predict(x_test, verbose=0)
y_pred = np.argmax(all_probs, axis=1)

def confusion_matrix_np(y_true, y_pred, num_classes=10):
    cm = np.zeros((num_classes, num_classes), dtype=np.int32)
    for t, p in zip(y_true, y_pred):
        cm[t, p] += 1
    return cm

cm = confusion_matrix_np(y_test, y_pred, num_classes=10)

plt.figure(figsize=(8,6))
plt.imshow(cm, cmap="viridis") # colored heatmap
plt.title("Confusion Matrix (Fashion-MNIST Test) - Colored")
plt.xlabel("Predicted Class")
plt.ylabel("True Class")
plt.xticks(range(10), class_names, rotation=35, ha="right")
plt.yticks(range(10), class_names)
plt.colorbar()

# write numbers on the heatmap for readability
thresh = cm.max() * 0.6
for i in range(10):
    for j in range(10):
        plt.text(j, i, cm[i, j],
            ha="center", va="center",
            color="white" if cm[i, j] > thresh else "black",
            fontsize=8)

```

```

plt.tight_layout()
plt.show()

# Print most common confusions (true -> predicted)
off_diag = cm.copy()
np.fill_diagonal(off_diag, 0)
pairs = []
for i in range(10):
    for j in range(10):
        if i != j and off_diag[i, j] > 0:
            pairs.append((off_diag[i, j], i, j))
pairs.sort(reverse=True)

print("\n==== MOST COMMON CONFUSIONS (true -> predicted) ===")
for c, t, p in pairs[:10]:
    print(f"- {class_names[t]} -> {class_names[p]}: {c} times")

# -----
# 9) Show misclassified examples
# -----
def show_misclassified(x_test_images, y_true, y_pred, probs, n=12):
    wrong_idx = np.where(y_true != y_pred)[0]
    if len(wrong_idx) == 0:
        print("No misclassifications found (rare).")
        return
    pick = np.random.choice(wrong_idx, size=min(n, len(wrong_idx)), replace=False)

    cols = 6
    rows = int(np.ceil(len(pick) / cols))
    plt.figure(figsize=(12, 2.2*rows))
    for i, k in enumerate(pick):
        plt.subplot(rows, cols, i+1)
        plt.imshow(x_test_images[k], cmap="gray")
        conf = np.max(probs[k])
        plt.title(f"T={class_names[y_true[k]]}\nP={class_names[y_pred[k]]}\n({conf*100:.1f}%)")
        plt.axis("off")
    plt.suptitle("Misclassified Test Examples (Fashion-MNIST)")
    plt.tight_layout()
    plt.show()

show_misclassified(x_test_img, y_test, y_pred, all_probs, n=12)

# -----
# 10) Random test cases (visual output + confidence)
# -----
def predict_test_cases(model, x_test_flat, x_test_images, y_test, n_cases=12):
    idx = np.random.choice(len(x_test_flat), n_cases, replace=False)
    probs = model.predict(x_test_flat[idx], verbose=0)
    preds = np.argmax(probs, axis=1)
    confs = np.max(probs, axis=1)

    cols = 6
    rows = int(np.ceil(n_cases / cols))
    plt.figure(figsize=(12, 2.2*rows))
    for i, k in enumerate(idx):
        plt.subplot(rows, cols, i+1)
        plt.imshow(x_test_images[k], cmap="gray")
        correct = (preds[i] == y_test[k])
        mark = "✓" if correct else "✗"
        plt.title(f"T={class_names[y_test[k]]}\nP={class_names[preds[i]]}\n({confs[i]*100:.1f}% {mark})")
        plt.axis("off")
    plt.suptitle("Random Test Cases: True vs Predicted + Confidence (Fashion-MNIST)")
    plt.tight_layout()
    plt.show()

    print("\nSample Predictions (index, true, pred, confidence):")
    for i, k in enumerate(idx):
        print(f"- idx={k:5d} | true={class_names[y_test[k]]:12s} | pred={class_names[preds[i]]:12s} | conf={confs[i]:.4f}")

predict_test_cases(model, x_test, x_test_img, y_test, n_cases=12)



# -----
# 11) Small detailed report
# -----
train_loss_last = history.history["loss"][-1]
val_loss_last = history.history["val_loss"][-1]
train_acc_last = history.history["accuracy"][-1]
val_acc_last = history.history["val_accuracy"][-1]



print("\n" + "="*60)
print("SMALL DETAILED REPORT (Task 2: Fashion-MNIST)")


```

```
print("=*60)
print("Exploration:")
print("- Loaded Fashion-MNIST: 60,000 train and 10,000 test grayscale images (28x28).")
print("- Printed shapes/pixel ranges, plotted class distributions, and showed sample images.")
print("\nPreprocessing:")
print("- Normalized pixel values to [0, 1].")
print("- Flattened each image from 28x28 to a 784-length vector.")
print("\nModel (Dense-only):")
print("- Dense(256, ReLU) + Dropout(0.3) + Dense(128, ReLU) + Dense(10, Softmax).")
print("\nTraining (10 epochs):")
print(f"- Final Train Loss: {train_loss_last:.4f} | Final Val Loss: {val_loss_last:.4f}")
print(f"- Final Train Acc : {train_acc_last:.4f} | Final Val Acc : {val_acc_last:.4f}")
print("\nTesting:")
print(f"- Test Loss: {test_loss:.4f} | Test Accuracy: {test_acc:.4f}")
print("\nError Analysis:")
print("- Built a colored confusion matrix (with counts in each cell).")
print("- Printed the most frequent confusions and visualized misclassified examples.")
print("=*60)
```

### Task 3 :







#### Preprocessing:

- Normalized pixel values to [0, 1].
- Flattened each image from 32x32x3 to 3072 features.

#### Model (Dense-only):

- Dense(512, ReLU) + Dropout(0.4) + Dense(256, ReLU) + Dropout(0.3) + Dense(10, Softmax).

#### Training (10 epochs):

- Final Train Acc: 0.3284 | Final Val Acc: 0.3956
- Final Train Loss: 1.8299 | Final Val Loss: 1.7847

#### Testing:

- Test Loss: 1.7659 | Test Accuracy: 0.4038

```

import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

tf.random.set_seed(42)
np.random.seed(42)

# -----
# 1) Load CIFAR-10
# -----
(x_train_img, y_train), (x_test_img, y_test) = tf.keras.datasets.cifar10.load_data()
y_train = y_train.squeeze() # (50000,1) -> (50000,)
y_test = y_test.squeeze() # (10000,1) -> (10000,)

class_names = [
    "airplane", "automobile", "bird", "cat", "deer",
    "dog", "frog", "horse", "ship", "truck"
]

print("==== DATASET OVERVIEW (CIFAR-10) ===")
print("Train images shape:", x_train_img.shape, " Train labels shape:", y_train.shape)
print("Test images shape:", x_test_img.shape, " Test labels shape:", y_test.shape)
print("Pixel range (train):", (x_train_img.min(), x_train_img.max()))
print("Pixel range (test):", (x_test_img.min(), x_test_img.max()))
print("Classes (0-9):", list(range(10)))
print("Class names:", class_names)

# -----
# 2) Explore dataset: distribution + sample images
# -----
def plot_class_distribution(labels, title):
    counts = np.bincount(labels, minlength=10)
    plt.figure(figsize=(10,3))
    plt.bar(range(10), counts)
    plt.xticks(range(10), class_names, rotation=20, ha="right")
    plt.xlabel("Class")
    plt.ylabel("Count")
    plt.title(title)
    plt.tight_layout()
    plt.show()
    return counts

train_counts = plot_class_distribution(y_train, "Train Class Distribution (CIFAR-10)")
test_counts = plot_class_distribution(y_test, "Test Class Distribution (CIFAR-10)")
print("Train class counts:", train_counts)
print("Test class counts:", test_counts)

def show_samples(images, labels, n=12, title="Random Samples"):
    idx = np.random.choice(len(images), n, replace=False)
    cols = 6
    rows = int(np.ceil(n / cols))
    plt.figure(figsize=(12, 2.4*rows))
    for i, k in enumerate(idx):
        plt.subplot(rows, cols, i+1)
        plt.imshow(images[k])
        plt.title(class_names[labels[k]])
        plt.axis("off")
    plt.suptitle(title)
    plt.tight_layout()
    plt.show()

show_samples(x_train_img, y_train, n=12, title="Random Training Samples (CIFAR-10)")

# -----
# 3) Preprocess: normalize + flatten
# -----
x_train = (x_train_img.astype("float32") / 255.0).reshape(-1, 32*32*3)
x_test = (x_test_img.astype("float32") / 255.0).reshape(-1, 32*32*3)

# -----
# 4) Build Dense-only model
#     (CIFAR-10 is harder for dense-only; keep it small-ish but capable)
# -----
model = tf.keras.Sequential([
    tf.keras.layers.Input(shape=(32*32*3,)),
    tf.keras.layers.Dense(512, activation="relu"),
    tf.keras.layers.Dropout(0.4),
    tf.keras.layers.Dense(256, activation="relu"),
    tf.keras.layers.Dropout(0.3),
    tf.keras.layers.Dense(10, activation="softmax")
])

model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

```

```

print("\n==== MODEL SUMMARY ===")
model.summary()

# -----
# 5) Train (10 epochs)
# -----
EPOCHS = 10
history = model.fit(
    x_train, y_train,
    epochs=EPOCHS,
    batch_size=128,
    validation_split=0.1,
    verbose=1
)

# -----
# 6) Plot training + validation accuracy curves (Task requirement)
# -----
def plot_accuracy_curves(hist):
    h = hist.history
    ep = range(1, len(h["accuracy"]) + 1)
    plt.figure(figsize=(7,4))
    plt.plot(ep, h["accuracy"], label="Train Accuracy")
    plt.plot(ep, h["val_accuracy"], label="Validation Accuracy")
    plt.xlabel("Epoch")
    plt.ylabel("Accuracy")
    plt.title("Accuracy Curves (CIFAR-10, Dense-only)")
    plt.legend()
    plt.tight_layout()
    plt.show()

plot_accuracy_curves(history)

# (Optional bonus: loss curves)
def plot_loss_curves(hist):
    h = hist.history
    ep = range(1, len(h["loss"]) + 1)
    plt.figure(figsize=(7,4))
    plt.plot(ep, h["loss"], label="Train Loss")
    plt.plot(ep, h["val_loss"], label="Validation Loss")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.title("Loss Curves (CIFAR-10, Dense-only)")
    plt.legend()
    plt.tight_layout()
    plt.show()

plot_loss_curves(history)

# -----
# 7) Test evaluation
# -----
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)
print("\n==== TEST PERFORMANCE ===")
print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_acc:.4f}")

# -----
# 8) Colored Confusion Matrix + Most common confusions
# -----
all_probs = model.predict(x_test, verbose=0)
y_pred = np.argmax(all_probs, axis=1)

def confusion_matrix_np(y_true, y_pred, num_classes=10):
    cm = np.zeros((num_classes, num_classes), dtype=np.int32)
    for t, p in zip(y_true, y_pred):
        cm[t, p] += 1
    return cm

cm = confusion_matrix_np(y_test, y_pred, num_classes=10)

plt.figure(figsize=(8,6))
plt.imshow(cm, cmap="plasma") # colored heatmap
plt.title("Confusion Matrix (CIFAR-10 Test) - Colored")
plt.xlabel("Predicted Class")
plt.ylabel("True Class")
plt.xticks(range(10), class_names, rotation=25, ha="right")
plt.yticks(range(10), class_names)
plt.colorbar()

# write counts on cells
thresh = cm.max() * 0.6
for i in range(10):
    for j in range(10):
        plt.text(j, i, cm[i, j],
                 ha="center", va="center",
                 color="white" if cm[i, j] > thresh else "black",
                 fontsize=8)

```

```

plt.tight_layout()
plt.show()

# show most frequent confusions
off_diag = cm.copy()
np.fill_diagonal(off_diag, 0)
pairs = []
for i in range(10):
    for j in range(10):
        if i != j and off_diag[i, j] > 0:
            pairs.append((off_diag[i, j], i, j))
pairs.sort(reverse=True)

print("\n==== MOST COMMON CONFUSIONS (true -> predicted) ===")
for c, t, p in pairs[:10]:
    print(f"- {class_names[t]} -> {class_names[p]}: {c} times")

# -----
# 9) Test cases (random predictions shown as images)
# -----
def predict_test_cases(x_test_images, y_true, y_pred, probs, n_cases=12):
    idx = np.random.choice(len(x_test_images), n_cases, replace=False)
    cols = 6
    rows = int(np.ceil(n_cases / cols))
    plt.figure(figsize=(12, 2.4*rows))
    for i, k in enumerate(idx):
        plt.subplot(rows, cols, i+1)
        plt.imshow(x_test_images[k])
        conf = np.max(probs[k])
        correct = (y_pred[k] == y_true[k])
        mark = "✓" if correct else "✗"
        plt.title(f"T={class_names[y_true[k]]}\nP={class_names[y_pred[k]]}\n{conf*100:.1f}% {mark}")
        plt.axis("off")
    plt.suptitle("Random Test Cases: True vs Predicted + Confidence (CIFAR-10)")
    plt.tight_layout()
    plt.show()

print("\nSample Predictions (index, true, pred, confidence):")
for k in idx:
    print(f"- idx={k:5d} | true={class_names[y_true[k]]}:10s} | pred={class_names[y_pred[k]]}:10s} | conf={np.max(probs[k]):.4f}")

predict_test_cases(x_test_img, y_test, y_pred, all_probs, n_cases=12)

# -----
# 10) Misclassified examples (visual)
# -----
def show_misclassified(x_test_images, y_true, y_pred, probs, n=12):
    wrong_idx = np.where(y_true != y_pred)[0]
    if len(wrong_idx) == 0:
        print("No misclassifications found (rare).")
        return
    pick = np.random.choice(wrong_idx, size=min(n, len(wrong_idx)), replace=False)

    cols = 6
    rows = int(np.ceil(len(pick) / cols))
    plt.figure(figsize=(12, 2.4*rows))
    for i, k in enumerate(pick):
        plt.subplot(rows, cols, i+1)
        plt.imshow(x_test_images[k])
        conf = np.max(probs[k])
        plt.title(f"T={class_names[y_true[k]]}\nP={class_names[y_pred[k]]}\n{conf*100:.1f}%")
        plt.axis("off")
    plt.suptitle("Misclassified Test Examples (CIFAR-10)")
    plt.tight_layout()
    plt.show()

show_misclassified(x_test_img, y_test, y_pred, all_probs, n=12)

# -----
# 11) Small detailed report
# -----
train_acc_last = history.history["accuracy"][-1]
val_acc_last = history.history["val_accuracy"][-1]
train_loss_last = history.history["loss"][-1]
val_loss_last = history.history["val_loss"][-1]

print("\n" + "="*60)
print("SMALL DETAILED REPORT (Task 3: CIFAR-10)")
print("=="*60)
print("Exploration:")
print("- Loaded CIFAR-10: 50,000 training images and 10,000 test images.")
print("- Each image is 32x32 with 3 color channels (RGB).")
print("- Displayed random sample images and plotted class distributions.")
print("\nPreprocessing:")
print("- Normalized pixel values to [0, 1].")
print("- Flattened each image from 32x32x3 to 3072 features.")
print("\nModel (Dense-only):")

```

```
print("- Dense(512, ReLU) + Dropout(0.4) + Dense(256, ReLU) + Dropout(0.3) + Dense(10, Softmax).")
print("\nTraining (10 epochs):")
print(f"- Final Train Acc: {train_acc_last:.4f} | Final Val Acc: {val_acc_last:.4f}")
print(f"- Final Train Loss: {train_loss_last:.4f} | Final Val Loss: {val_loss_last:.4f}")
print("\nTesting:")
print(f"- Test Loss: {test_loss:.4f} | Test Accuracy: {test_acc:.4f}")
print("\nNotes:")
print("- CIFAR-10 is harder than MNIST/Fashion-MNIST; Dense-only models usually perform worse than CNNs.")
print("- Confusion matrix and misclassified examples help identify which classes are most frequently mixed up.")
print("=*60)
```

# Conclusion

In these three tasks, we built and tested fully connected (dense) neural networks on three different image classification datasets: MNIST, Fashion-MNIST, and CIFAR-10.

The goal was to understand how dense neural networks work, how data is prepared, and how model performance changes with dataset complexity.

## Task 1: MNIST (Handwritten Digits)

- Loaded the MNIST dataset of handwritten digits (0–9).
- Normalized the pixel values and flattened each image.
- Built a small neural network using only dense layers.
- Trained the model for 10 epochs.
- Evaluated the model using accuracy, loss curves, confusion matrix, and test examples.
- The model achieved **very high accuracy**.
- Training and validation curves showed stable learning.
- The confusion matrix showed very few errors.
- This task proved that dense networks work **very well** for simple image data.

## Conclusion

MNIST is an easy dataset, and a simple fully connected neural network is enough to classify digits accurately.

## Task 2: Fashion-MNIST (Clothing Images)

- Explored clothing image classes (shirts, shoes, bags, etc.).
- Normalized and flattened the images.
- Built a dense neural network with dropout for regularization.
- Trained and tested the model.
- Visualized accuracy, loss curves, confusion matrix, and misclassified samples.
- Accuracy was lower than MNIST, but still reasonable.
- The confusion matrix showed that similar clothes (e.g., shirt vs coat) are often confused.
- Misclassified images helped us understand where the model struggles.

## Conclusion

Fashion-MNIST is more complex than MNIST. Dense networks can still work, but they start to struggle with visual similarity between classes.

### **Task 3: CIFAR-10 (Color Object Images)**

- Worked with color images (airplanes, cars, animals, etc.).
- Normalized and flattened RGB images.
- Built a larger dense network to handle higher complexity.
- Trained the model and plotted training and validation accuracy.
- Analyzed results using confusion matrix and test predictions.
- Accuracy was significantly lower than the previous tasks.
- The model confused similar objects (e.g., cats vs dogs, trucks vs cars).
- This showed the limitations of dense-only networks for complex image data.

### **Conclusion**

CIFAR-10 is a hard dataset for fully connected networks. This task clearly shows why Convolutional Neural Networks (CNNs) are better for image classification.